

Assignment #4

Student name: *Hardik Gupta*

Course: *Computer Vision (CSCI 5561)* – Professor: *Dr. Volkan Isler*
Due date: *November 20th, 2023*

Tiny Image kNN Classification

```
def get_tiny_image(img, tiny_size):
```

```
    ...  
    return feature
```

Input: *img* is an gray scale image, *tiny_size* = (*w*, *h*) is the tiny image size.

Output: *feature* is a vector of size *wh1* representing the tiny image.

```
def predict_kNN(feature_train, label_train, feature_test, n_neighbors):
```

```
    ...  
    return pred_test
```

Input: *feature_train* is a $n_{tr} \times d$ matrix where n_{tr} is the number of training data samples and d is the dimension of image feature, e.g., 256 for 16×16 tiny image representation. Each row is the image feature. *label_train* $\in [1, 15]$ is a n_{tr} vector that specifies the label of the training data. *feature_test* is a $n_{te} \times d$ matrix that contains the testing features where n_{te} is the number of testing data samples. *n_neighbors* is the number of neighbors for label prediction.

Output: *pred_test* is a n_{te} vector that specifies the predicted label for the testing data.

```
def compute_confusion_matrix_and_accuracy(pred_test, label_test, n_classes):
```

```
    ...  
    return confusion, accuracy
```


Input: *pred_test* is a n_{te} vector that specifies the predicted label for the testing data. *label_test* $\in [1, 15]$ is a n_{te} vector that specifies the label of the test data. *n_classes* is the number of categories, i.e. 15 in this case.

Output: *confusion* is a 15×15 confusion matrix and *accuracy* is the prediction accuracy on testing data.

Solution. Resizing of the image in *tiny_size*. Then vectorization and and normalization of the image to have 0 mean and unit length.

We create an instance of the *KNeighborsClassifier* and fitting the label and features of the samples and produce predictions on test samples.

Then we produce *confusion_matrix* and accuracy score using the predictions.



Bag-of-words Visual Vocabulary

```
def compute_dsift(img, stride, size):
```

```
    ...  
    return dsift
```

Input: *img* is a gray scale image. *stride* and *size* are both integers controls locations on image to compute SIFT features and diameter of the meaningful keypoint neighborhood.

Output: *dsift* is a collection of SIFT features whose size is $n \times 128$. *n* is total number of locations to compute SIFT features on *img*.

```
def build_visual_dictionary(features, d_size):
```

```
    ...  
    return vocab
```

Input: features is a collection of SIFT features of all training images ($N \times 128$) and *d_size* is the size of the dictionary (the number of visual words).

Output: vocab lists the quantized visual words whose size is $d_size \times 128$.

```
def compute_bow(dsift, vocab):
```

```
    ...  
    return bow_feature
```


Input: *dsift* is a set of SIFT features for one image, and vocab is visual dictionary.

Output: *bow_feature* is the bag-of-words feature vector whose size is *d_size*.

Solution. **compute_dsift** function computes dense SIFT (Scale-Invariant Feature Transform) features for an input image with a specified stride and size.

build_visual_dictionary function builds a visual dictionary using K-means clustering (KMeans). It takes a set of features (features) and the desired dictionary size (dict_size). The resulting dictionary is returned.

compute_bow function computes a Bag of Visual Words (BoW) feature representation for a set of dense SIFT features (dsift) using a pre-built visual dictionary (vocab).



BoW+SVM

```
def predict_svm(feature_train, label_train, feature_test, n_classes):
    ...
    return pred_test
```

Input: *feature_train* is a $n_{tr} \times d$ matrix where n_{tr} is the number of training data samples and d is the dimension of image feature. Each row is the image feature. *label_train* $\in [1, 15]$ is a n_{tr} vector that specifies the label of the training data. *feature_test* is a $n_{te} \times d$ matrix that contains the testing features where n_{te} is the number of testing data samples. *n_classes* is the number of categories, i.e., 15 in this case.

Output: *pred_test* is a n_{tr} vector that specifies the predicted label for the testing data.

Solution. `predict_svm` function performs Support Vector Machine (SVM) classification. It takes training features (*feature_train*) and labels (*label_train*), as well as test features (*feature_test*) and the number of classes (*n_classes*). It returns the predicted labels for the test features using an SVM classifier.

