

## Assignment #5

Student name: *Hardik Gupta*

---

Course: *Computer Vision (CSCI 5561)* – Professor: *Dr. Volkan Isler*  
Due date: *December 4th, 2023*

### Multi Layer Perceptron

```
def get_mini_batch(im_train, label_train, batch_size)
```

```
...
```

```
    return mini_batch_x, mini_batch_y
```

**Input:**  $\text{im\_train} \in \mathbb{R}^{196 \times N}$  and  $\text{label\_train} \in \mathbb{R}_1 \times N$  are a set of vectorized images and corresponding labels, and  $\text{batch\_size}$  is the size of the mini-batch for stochastic gradient descent.

**Output:**  $\text{mini\_batch\_x}$  and  $\text{mini\_batch\_y}$  are python lists that contain a set of batches (images and labels, respectively). Each batch of images is a matrix with size  $196 \times \text{batch\_size}$ , and each batch of labels is a matrix with size  $10 \times \text{batch\_size}$  (one-hot encoding). Note that the number of images in the last batch may be smaller than  $\text{batch\_size}$ .

```
def fc(x, w, b)
```

```
...
```

```
    return y
```

**Input:**  $x \in \mathbb{R}^{m \times 1}$  is the input to the fully connected layer, and  $w \in \mathbb{R}^{n \times m}$  and  $b \in \mathbb{R}^{n \times 1}$  are the weights and bias.

**Output:**  $y \in \mathbb{R}^{n \times 1}$  is the output of the linear transform (fully connected layer).

```
def fc_backward(dl_dy, x, w, b)
```

```
...
```

```
    return dl_dx, dl_dw, dl_db
```

**Input:**  $\text{dl\_dy} \in \mathbb{R}^{n \times 1}$  is the loss derivative with respect to the output  $y$  of fully connected layer.

**Output:**  $\text{dl\_dx} \in \mathbb{R}^{m \times 1}$  is the loss derivative with respect the input  $x$ ,  $\text{dl\_dw} \in \mathbb{R}^{n \times m}$  is the loss derivative with respect to the weights, and  $\text{dl\_db} \in \mathbb{R}^{n \times 1}$  is the loss derivative with respect to the bias.

```
def relu(x)
    ...
    return y
```

**Input:**  $x$  is a general tensor, matrix, or vector.

**Output:**  $y$  is the output of the Rectified Linear Unit (ReLU) with the same shape as input.

```
def relu_backward(dl_dy, x)
    ...
    return dl_dx
```

**Input:**  $dl\_dy$  is the loss derivative with respect to the output  $y$  of ReLU layer. It has the same shape as  $y$  (it can be a tensor, matrix, or vector).

**Output:**  $dl\_dx$  is the loss derivative with respect to the input  $x$ . It has the same shape as  $dl\_dy$ .

```
def loss_cross_entropy_softmax(x, y)
    ...
    return l, dl_dx
```

**Input:**  $x \in \mathbb{R}^{m \times 1}$  is the input to the softmax, and  $y \in 0, 1^{m \times 1}$  is the ground truth label.

**Output:**  $l \in \mathbb{R}$  is the loss, and  $dl\_dx \in \mathbb{R}^{m \times 1}$  is the loss derivative with respect to  $x$ .

**Solution.** `get_mini_batch` iterates through the batches, creating mini-batches of input features (`mini_batch_x`) and their corresponding one-hot encoded labels (`mini_batch_y`). The one-hot encoding is applied to the labels, converting them into a binary matrix representation. Overall, this function sets up the data in a randomized and batched format, ready for training.

**fc** linear transform of  $x$ , i.e.,  $y = wx + b$

**fc\_backward** provides mathematical calculation for each:


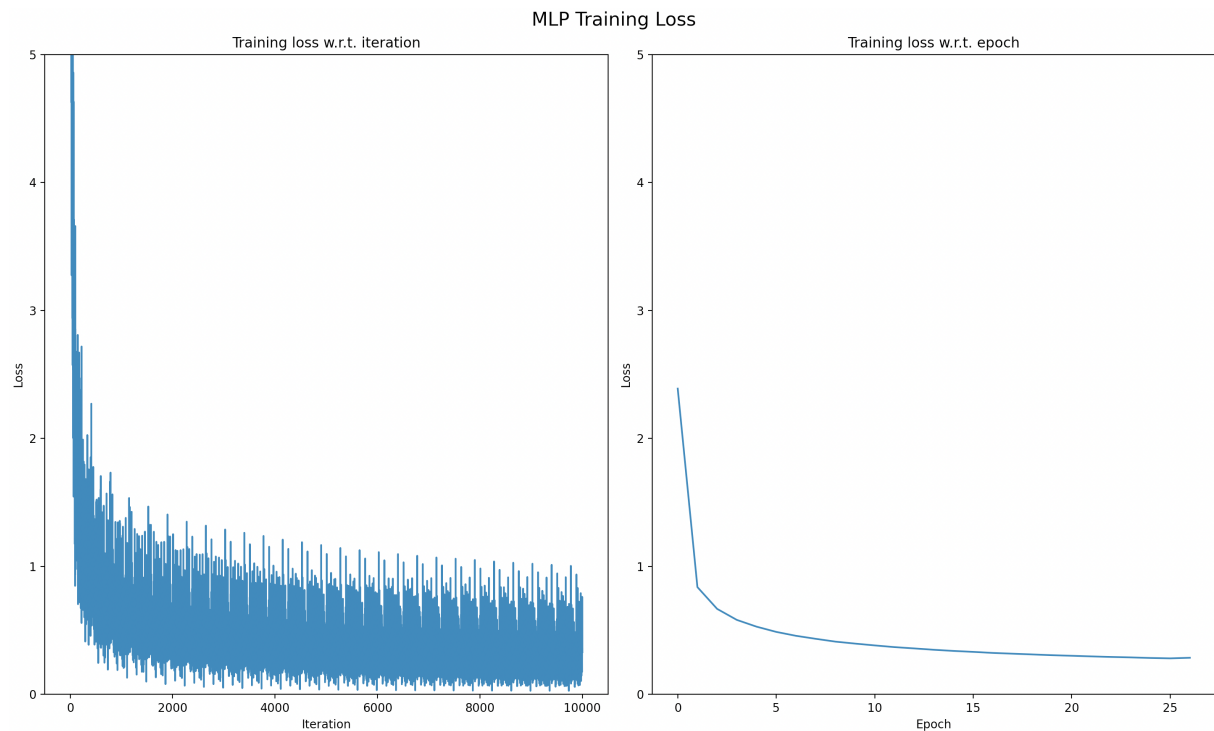
1.  $\frac{dl}{dx} = \frac{dl}{dy} \cdot \frac{dy}{dx}$ ; where  $\frac{dy}{dx} = w$
2.  $\frac{dl}{dw} = \frac{dl}{dy} \cdot \frac{dy}{dw}$ ; where  $\frac{dy}{dw} = x$
3.  $\frac{dl}{db} = \frac{dl}{dy} \cdot \frac{dy}{db}$ ; where  $\frac{dy}{db} = 1$

**relu** provides the maximum of  $(0, x)$

**relu\_backward** provides the  $\frac{dy}{dx}$ ; where  $y = \text{ReLU}$ .  $\frac{dl}{dx} = \frac{dl}{dy} \frac{dy}{dx}$ ,

$\frac{dy}{dx} = 1$  where  $x > 0$ ; else 0

**loss\_cross\_entropy** is calculated using the ground truth label  $y$  and the softmax output  $y\_hat$ . It measures the dissimilarity between the predicted distribution and the true distribution. The *softmax* function is applied element-wise to the input vector  $x$ . It exponentiates each element, subtracts the maximum value to improve numerical stability, and then normalizes the result to obtain a probability distribution. The derivative of the loss with respect to the input  $x$  is computed as the difference between the softmax output  $y\_hat$  and the ground truth label  $y$ .



## Convolutional Neural Network

```
def conv(x, w_conv, b_conv)
```

```
    ...
    return y
```

**Input:**  $x \in \mathbb{R}^{H \times W \times C_1}$  is an input to the convolutional operation,  $w_{\text{conv}} \in \mathbb{R}^{h \times w \times C_1 \times C_2}$  and  $b_{\text{conv}} \in \mathbb{R}^{C_2 \times 1}$  are weights and bias of the convolutional operation.

**Output:**  $y \in \mathbb{R}^{H \times W \times C_2}$  is the output of the convolutional operation. Note that to get the same size with the input, you may pad zero at the boundary of the input image.

```
def conv_backward(dl_dy, x, w_conv, b_conv)
```

```
    ...
    return dl_dw, dl_db
```

**Input:**  $dl\_dy \in \mathbb{R}^{H \times W \times C_2}$  is the loss derivative with respect to the output of the convolutional layer. The rest inputs are the same as defined in function *conv*.

**Output:**  $dl\_dw \in \mathbb{R}^{h \times w \times C_1 \times C_2}$  and  $dl\_db \in \mathbb{R}^{C_2 \times 1}$  are the loss derivatives with respect to convolutional weights  $w_{\text{conv}}$  and bias  $b_{\text{conv}}$ , respectively.

```
def pool2x2(x)
```

```
    ...
    return y
```

**Input:**  $x \in \mathbb{R}^{H \times W \times C}$  is a general tensor or matrix.

**Output:**  $y \in \mathbb{R}^{\frac{H}{2} \times \frac{W}{2} \times C}$  is the output of the  $2 \times 2$  max-pooling operation with stride 2.

```
def pool2x2_backward(dl_dy, x)
```

```
    ...
    return dl_dx
```

**Input:**  $dl\_dy$  is the loss derivative with respect to the output  $y$ .  $x \in \mathbb{R}^{H \times W \times C}$  is the input to the pooling layer.

**Output:**  $dl\_dx \in \mathbb{R}^{H \times W \times C}$  is the loss derivative with respect to the input  $x$ .

**Solution.** *conv* extracts the dimensions of the input tensor and the convolutional filter. It then calculates the padding required to maintain spatial dimensions during convolution. The input tensor is padded using NumPy's *np.pad* function to handle border effects. For each spatial position, the function extracts the corresponding region from the padded input tensor, performs element-wise multiplication with the convolutional filter, and sums the results along all dimensions except the channel dimension. The bias term is added to the sum. The final output tensor  $y$  represents the result of applying the convolution operation to the input tensor  $x$  with the specified convolutional filter  $w_{\text{conv}}$  and bias term  $b_{\text{conv}}$ .

*conv\_backward* function performs the backpropagation step for a 2D convolution operation. It calculates and accumulates the gradients with respect to the filter weights and biases, allowing for the subsequent optimization of these parameters during the

training of a neural network.

The **pool2x2** function implements a  $2 \times 2$  max pooling operation on a given input tensor  $x$ . The function iterates over the input tensor and, for each  $2 \times 2$  patch in each channel, extracts the maximum value. The result is stored in an output tensor  $y$ , effectively reducing the spatial dimensions by half in both height and width. The function employs a stride of 2 and is designed for 2D convolutional neural networks (CNNs) to down-sample feature maps, retaining essential information while reducing computational complexity.

**pool2x2\_backward** function calculates the gradient of the loss with respect to the input tensor in the context of a  $2 \times 2$  max pooling operation. It iterates through the output gradient tensor ( $dl\_dy$ ), identifies the corresponding  $2 \times 2$  block in the input tensor  $x$ , and assigns the gradient values to the position of the maximum element. This facilitates the backpropagation process, crucial for updating model parameters during training in convolutional neural networks.

